

2018年5月22日 19:07

[illegible]

可以看到write其实调用的是个函数指针，这就是vfs的层核心结构，通过file_operations挂接到对应的不通文件系统中：

```

struct file {
    union {
        struct llist_node    fu_llist;
        struct rcu_head      fu_rcuhead;
    } f_u;
    struct path              f_path;
    struct inode              *f_inode; /* cached value */
    const struct file_operations *f_op;

    /*

```

例如ext4文件系统创建inode时初始化函数指针:

```

const struct file_operations ext4_file_operations = {
    .llseek      = ext4_llseek,
    .read_iter   = generic_file_read_iter,
    .write_iter  = ext4_file_write_iter,
    .unlocked_ioctl = ext4_ioctl,
#ifdef CONFIG_COMPAT
    .compat_ioctl = ext4_compat_ioctl,
#endif
    .mmap        = ext4_file_mmap,
    .open        = ext4_file_open,
    .release     = ext4_release_file,
    .fsync       = ext4_sync_file,
    .get_unmapped_area = thp_get_unmapped_area,
    .splice_read = generic_file_splice_read,
    .splice_write = iter_file_splice_write,
    .fallocate   = ext4_fallocate,
};

```

```

static int ext4_create(struct inode *dir, struct dentry *dentry, umode_t mode,
                      bool excl)
{
    handle_t *handle;
    struct inode *inode;
    int err, credits, retries = 0;

    err = dqquot_initialize(dir);
    if (err)
        return err;

    credits = (EXT4_DATA_TRANS_BLOCKS(dir->i_sb) +
               EXT4_INDEX_EXTRA_TRANS_BLOCKS + 3);
retry:
    inode = ext4_new_inode_start_handle(dir, mode, &dentry->d_name, 0,
                                       NULL, EXT4_HT_DIR, credits);
    handle = ext4_journal_current_handle();
    err = PTR_ERR(inode);
    if (!IS_ERR(inode)) {
        inode->i_op = &ext4_file_inode_operations;
        inode->i_fop = &ext4_file_operations;
        ext4_set_aops(inode);
        err = ext4_add_nondir(handle, dentry, inode);
        if (!err && IS_DIRSYNC(dir))
            ext4_handle_sync(handle);
    }
    if (handle)
        ext4_journal_stop(handle);
    if (err == -ENOSPC && ext4_should_retry_alloc(dir->i_sb, &retries))
        goto ↑retry;
    return err;
} ? end ext4_create ?

```

基本关系沥青了我们开始分析ext4写文件的入口函数: ext4_file_write_iter

```

static ssize_t
ext4_file_write_iter(struct kiocb *iocb, struct iov_iter *from)
{
    struct inode *inode = file_inode(iocb->ki_filp); //找到文件对应的inode
    int o_direct = iocb->ki_flags & IOC_B_DIRECT; //判断是否是direct模式
    int unaligned_aio = 0;
    int overwrite = 0;
    ssize_t ret;

    inode_lock(inode);
    ret = generic_write_checks(iocb, from);
    if (ret <= 0)
        goto out;

    /*
     * Unaligned direct AIO must be serialized among each other as zeroing
     * of partial blocks of two competing unaligned AIOs can result in data
     * corruption.
     */
    if (o_direct && ext4_test_inode_flag(inode, EXT4_INODE_EXTENTS) &&
        !is_sync_kiocb(iocb) &&
        ext4_unaligned_aio(inode, from, iocb->ki_pos)) {
        unaligned_aio = 1;
        ext4_unwritten_wait(inode);
    }

    /*
     * If we have encountered a bitmap-format file, the size limit
     * is smaller than s_maxbytes, which is for extent-mapped files. //处理非extent文件
     */
    if (!!(ext4_test_inode_flag(inode, EXT4_INODE_EXTENTS))) {
        struct ext4_sb_info *sbi = EXT4_SB(inode->i_sb);

        if (iocb->ki_pos >= sbi->s_bitmap_maxbytes) {
            ret = -EFBIG;
            goto out;
        }
        iov_iter_truncate(from, sbi->s_bitmap_maxbytes - iocb->ki_pos);
    }

    iocb->private = &overwrite;
    if (o_direct) {
        size_t length = iov_iter_count(from);
        loff_t pos = iocb->ki_pos;

```

```

/* check whether we do a DIO overwrite or not */
if (ext4_should_dioread_nolock(inode) && !unaligned_aio &&
    pos + length <= i_size_read(inode)) {
    struct ext4_map_blocks map;
    unsigned int blkbits = inode->i_blkbits;
    int err, len;

    map.m_lblk = pos >> blkbits;
    map.m_len = EXT4_MAX_BLOCKS(length, pos, blkbits);
    len = map.m_len;

    err = ext4_map_blocks(NULL, inode, &map, 0);
    /*
     * 'err==len' means that all of blocks has
     * been preallocated no matter they are
     * initialized or not. For excluding
     * unwritten extents, we need to check
     * m_flags. There are two conditions that
     * indicate for initialized extents. 1) If we
     * hit extent cache, EXT4_MAP_MAPPED flag is
     * returned; 2) If we do a real lookup,
     * non-flags are returned. So we should check
     * these two conditions.
     */
    if (err == len && (map.m_flags & EXT4_MAP_MAPPED))
        overwrite = 1;
}

ret = __generic_file_write_iter(iocb, from);
inode_unlock(inode);

if (ret > 0)
    ret = generic_write_sync(iocb, ret);

return ret;

out:
inode_unlock(inode);
return ret;
}

```

可以看到申请block的入口是ext4_map_blocks, 真正写文件的入口是__generic_file_write_iter

```

/**
 * __generic_file_write_iter - write data to a file
 * @iocb:      IO state structure (file, offset, etc.)
 * @from:      iov_iter with data to write
 *
 * This function does all the work needed for actually writing data to a
 * file. It does all basic checks, removes SUID from the file, updates
 * modification times and calls proper subroutines depending on whether we
 * do direct IO or a standard buffered write.
 *
 * It expects i_mutex to be grabbed unless we work on a block device or similar
 * object which does not need locking at all.
 *
 * This function does *not* take care of syncing data in case of O_SYNC write.
 * A caller has to handle it. This is mainly due to the fact that we want to
 * avoid syncing under i_mutex.
 */
ssize_t __generic_file_write_iter(struct kiocb *iocb, struct iov_iter *from)
{
    struct file *file = iocb->ki_filp;
    struct address_space *mapping = file->f_mapping;
    struct inode *inode = mapping->host;
    ssize_t written = 0;
    ssize_t err;
    ssize_t status;

    /* We can write back this queue in page reclaim */
    current->backing_dev_info = inode_to_bdi(inode);
    err = file_remove_privs(file);
    if (err)
        goto out;

    err = file_update_time(file); //更新元数据信息中的时间信息
    if (err)
        goto out;

    if (iocb->ki_flags & IOCB_DIRECT) { //处理direct io
        loff_t pos, endbyte;

        written = generic_file_direct_write(iocb, from); //这里调用ext4_direct_IO
        /*
         * If the write stopped short of completing, fall back to
         * buffered writes. Some filesystems do this for writes to
         * holes, for example. For DAX files, a buffered write will
         * not succeed (even if it did, DAX does not handle dirty
         * page-cache pages correctly).
         */
        if (written < 0 || !iov_iter_count(from) || IS_DAX(inode))
            goto out;

        status = generic_perform_write(file, from, pos = iocb->ki_pos);
        /*
         * If generic_perform_write() returned a synchronous error
         * then we want to return the number of bytes which were
         * direct-written, or the error code if that was zero. Note
         * that this differs from normal direct-io semantics, which
         * will return -EFOO even if some bytes were written.
         */
        if (unlikely(status < 0)) {
            err = status;
            goto out;
        }
        /*
         * We need to ensure that the page cache pages are written to
         * disk and invalidated to preserve the expected O_DIRECT
         * semantics.
         */
    }
}

```

```

endbyte = pos + status - 1;
err = filemap_write_and_wait_range(mapping, pos, endbyte);
if (err == 0) {
    iocb->ki_pos = endbyte + 1;
    written += status;
    invalidate_mapping_pages(mapping,
                             pos >> PAGE_SHIFT,
                             endbyte >> PAGE_SHIFT);
} else {
    /*
     * We don't know how much we wrote, so just return
     * the number of bytes which were direct-written
     */
}
} else {
    //处理非direct io写, 调用write_begin和write_end
    written = generic_perform_write(file, from, iocb->ki_pos);
    if (likely(written > 0))
        iocb->ki_pos += written;
}
}
out:
current->backing_dev_info = NULL;
return written ? written : err;
}

```

以下是不通过mount选项下ext4写实现的函数实现

```

03618: static const struct address_space_operations ext4_aops = {
03619:     .readpage           = ext4_readpage,
03620:     .readpages          = ext4_readpages,
03621:     .writepage          = ext4_writepage,
03622:     .writepages         = ext4_writepages,
03623:     .write_begin        = ext4_write_begin,
03624:     .write_end          = ext4_write_end,
03625:     .bmap               = ext4_bmap,
03626:     .invalidatepage     = ext4_invalidatepage,
03627:     .releasepage        = ext4_releasepage,
03628:     .direct_IO          = ext4_direct_IO,
03629:     .migratepage        = buffer_migrate_page,
03630:     .is_partially_uptodate = block_is_partially_uptodate,
03631:     .error_remove_page  = generic_error_remove_page,
03632: };
03633:
03634: static const struct address_space_operations ext4_journalled_aops = {
03635:     .readpage           = ext4_readpage,
03636:     .readpages          = ext4_readpages,
03637:     .writepage          = ext4_writepage,
03638:     .writepages         = ext4_writepages,
03639:     .write_begin        = ext4_write_begin,
03640:     .write_end          = ext4_journalled_write_end,
03641:     .set_page_dirty     = ext4_journalled_set_page_dirty,
03642:     .bmap               = ext4_bmap,
03643:     .invalidatepage     = ext4_journalled_invalidatepage,
03644:     .releasepage        = ext4_releasepage,
03645:     .direct_IO          = ext4_direct_IO,
03646:     .is_partially_uptodate = block_is_partially_uptodate,
03647:     .error_remove_page  = generic_error_remove_page,
03648: };
03649:
03650: static const struct address_space_operations ext4_da_aops = {
03651:     .readpage           = ext4_readpage,
03652:     .readpages          = ext4_readpages,
03653:     .writepage          = ext4_writepage,
03654:     .writepages         = ext4_writepages,
03655:     .write_begin        = ext4_da_write_begin,
03656:     .write_end          = ext4_da_write_end,
03657:     .bmap               = ext4_bmap,
03658:     .invalidatepage     = ext4_da_invalidatepage,
03659:     .releasepage        = ext4_releasepage,
03660:     .direct_IO          = ext4_direct_IO,
03661:     .migratepage        = buffer_migrate_page,
03662:     .is_partially_uptodate = block_is_partially_uptodate,

```

```

Ext4 direct io:
static ssize_t ext4_direct_IO(struct kiocb *iocb, struct iov_iter *iter)
{
    struct file *file = iocb->ki_filp;
    struct inode *inode = file->f_mapping->host;
    size_t count = iov_iter_count(iter);
    loff_t offset = iocb->ki_pos;
    ssize_t ret;

#ifdef CONFIG_EXT4_FS_ENCRYPTION
    if (ext4_encrypted_inode(inode) && S_ISREG(inode->i_mode))
        return 0;
#endif

    /*
     * If we are doing data journaling we don't support O_DIRECT
     */
    if (ext4_should_journal_data(inode))
        return 0;

    /* Let buffer I/O handle the inline data case. */
    if (ext4_has_inline_data(inode))
        return 0;

    trace_ext4_direct_IO_enter(inode, offset, count, iov_iter_rw(iter));
    if (iov_iter_rw(iter) == READ)
        ret = ext4_direct_IO_read(iocb, iter); //direct read
    else
        ret = ext4_direct_IO_write(iocb, iter); //direct write
    trace_ext4_direct_IO_exit(inode, offset, count, iov_iter_rw(iter), ret);
    return ret;
}

```

分析direct写:

```

/*
 * Handling of direct IO writes.
 *
 * For ext4 extent files, ext4 will do direct-io write even to holes,
 * preallocated extents, and those write extend the file, no need to
 * fall back to buffered IO.

```

```

*
* For holes, we fallocate those blocks, mark them as unwritten
* If those blocks were preallocated, we mark sure they are split, but
* still keep the range to write as unwritten.
*
* The unwritten extents will be converted to written when DIO is completed.
* For async direct IO, since the IO may still pending when return, we
* set up an end_io call back function, which will do the conversion
* when async direct IO completed.
*
* If the O_DIRECT write will extend the file then add this inode to the
* orphan list. So recovery will truncate it back to the original size
* if the machine crashes during the write.
*
*/
static ssize_t ext4_direct_IO_write(struct kiocb *iocb, struct iov_iter *iter)
{
    struct file *file = iocb->ki_filp;
    struct inode *inode = file->f_mapping->host;
    struct ext4_inode_info *ei = EXT4_I(inode);
    ssize_t ret;
    loff_t offset = iocb->ki_pos;
    size_t count = iov_iter_count(iter);
    int overwrite = 0;
    get_block_t *get_block_func = NULL;
    int dio_flags = 0;
    loff_t final_size = offset + count;
    int orphan = 0;
    handle_t *handle;

    if (final_size > inode->i_size) {
        /* Credits for sb + inode write */
        handle = ext4_journal_start(inode, EXT4_HT_INODE, 2);
        if (IS_ERR(handle)) {
            ret = PTR_ERR(handle);
            goto out;
        }
        ret = ext4_orphan_add(handle, inode);
        if (ret) {
            ext4_journal_stop(handle);
            goto out;
        }
        orphan = 1;
        ei->i_disksize = inode->i_size;
        ext4_journal_stop(handle);
    }

    BUG_ON(iocb->private == NULL);

    /*
     * Make all waiters for direct IO properly wait also for extent
     * conversion. This also disallows race between truncate() and
     * overwrite DIO as i_dio_count needs to be incremented under i_mutex.
     */
    inode_dio_begin(inode);

    /* If we do a overwrite dio, i_mutex locking can be released */
    overwrite = *((int *)iocb->private);

    if (overwrite)
        inode_unlock(inode);

    /*
     * For extent mapped files we could direct write to holes and fallocate.
     *
     * Allocated blocks to fill the hole are marked as unwritten to prevent
     * parallel buffered read to expose the stale data before DIO complete
     * the data IO.
     *
     * As to previously fallocated extents, ext4 get_block will just simply
     * mark the buffer mapped but still keep the extents unwritten.
     *
     * For non AIO case, we will convert those unwritten extents to written
     * after return back from blockdev_direct_IO. That way we save us from
     * allocating io_end structure and also the overhead of offloading
     * the extent conversion to a workqueue.
     *
     * For async DIO, the conversion needs to be deferred when the
     * IO is completed. The ext4 end_io callback function will be
     * called to take care of the conversion work. Here for async
     * case, we allocate an io_end structure to hook to the iocb.
     */
    iocb->private = NULL;
    if (overwrite)
        get_block_func = ext4_dio_get_block_overwrite; // 几种获取block num的函数，核心是
        ext4_map_blocks
    else if (IS_DAX(inode)) {
        /*
         * We can avoid zeroing for aligned DAX writes beyond EOF. Other
         * writes need zeroing either because they can race with page
         * faults or because they use partial blocks.
         */
        if (round_down(offset, 1 << inode->i_blkbits) >= inode->i_size &&
            ext4_aligned_io(inode, offset, count))
            get_block_func = ext4_dio_get_block;
        else
            get_block_func = ext4_dax_get_block;
        dio_flags = DIO_LOCKING;
    } else if (!ext4_test_inode_flag(inode, EXT4_INODE_EXTENTS) ||
        round_down(offset, 1 << inode->i_blkbits) >= inode->i_size) {
        get_block_func = ext4_dio_get_block;
        dio_flags = DIO_LOCKING | DIO_SKIP_HOLES;
    } else if (is_sync_kiocb(iocb)) {
        get_block_func = ext4_dio_get_block_unwritten_sync;
        dio_flags = DIO_LOCKING;
    } else {
        get_block_func = ext4_dio_get_block_unwritten_async;
        dio_flags = DIO_LOCKING;
    }
}

#ifdef CONFIG_EXT4_FS_ENCRYPTION
    BUG_ON(ext4_encrypted_inode(inode) && S_ISREG(inode->i_mode));
#endif

```

```

if (IS_DAX(inode)) {
    ret = dax_do_io(iocb, inode, iter, get_block_func,
        ext4_end_io_dio, dio_flags);
} else
    ret = __blockdev_direct_IO(iocb, inode, //提交到block layer
        inode->i_sb->s_bdev, iter,
        get_block_func,
        ext4_end_io_dio, NULL, dio_flags);

if (ret > 0 && !overwrite && ext4_test_inode_state(inode,
    EXT4_STATE_DIO_UNWRITTEN)) {
    int err;
    /*
     * for non AIO case, since the IO is already
     * completed, we could do the conversion right here
     */
    err = ext4_convert_unwritten_extents(NULL, inode,
        offset, ret);
    if (err < 0)
        ret = err;
    ext4_clear_inode_state(inode, EXT4_STATE_DIO_UNWRITTEN);
}

inode_dio_end(inode);
/* take i_mutex locking again if we do a overwrite dio */
if (overwrite)
    inode_lock(inode);

if (ret < 0 && final_size > inode->i_size)
    ext4_truncate_failed_write(inode);

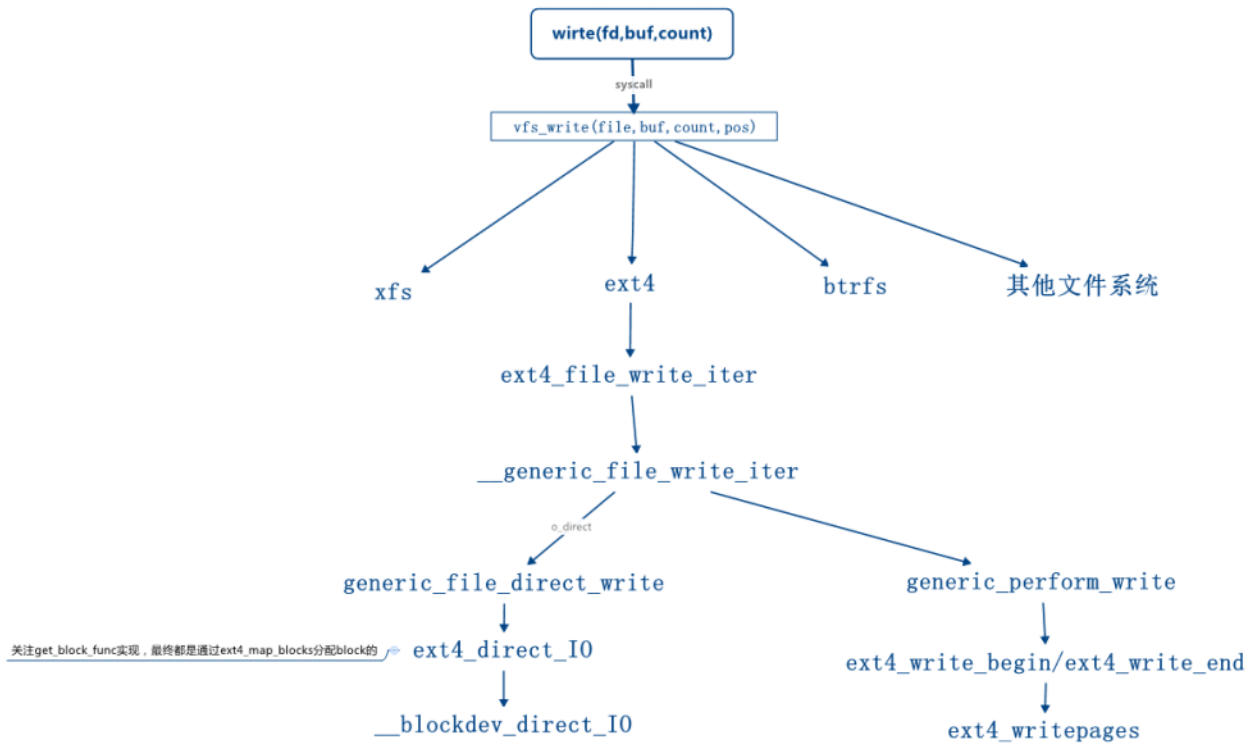
/* Handle extending of i_size after direct IO write */
if (orphan) {
    int err;

    /* Credits for sb + inode write */
    handle = ext4_journal_start(inode, EXT4_HT_INODE, 2);
    if (IS_ERR(handle)) {
        /* This is really bad luck. We've written the data
         * but cannot extend i_size. Bail out and pretend
         * the write failed... */
        ret = PTR_ERR(handle);
        if (inode->i_nlink)
            ext4_orphan_del(NULL, inode);

        goto out;
    }
    if (inode->i_nlink)
        ext4_orphan_del(handle, inode);
    if (ret > 0) {
        loff_t end = offset + ret;
        if (end > inode->i_size) {
            ei->i_disksize = end;
            i_size_write(inode, end);
            /*
             * We're going to return a positive `ret'
             * here due to non-zero-length I/O, so there's
             * no way of reporting error returns from
             * ext4_mark_inode_dirty() to userspace. So
             * ignore it.
             */
            ext4_mark_inode_dirty(handle, inode);
        }
    }
    err = ext4_journal_stop(handle);
    if (ret == 0)
        ret = err;
}
out:
    return ret;
}

```

最后梳理一下整个的调用流程：



其中get_block_func要实现的就是文件的offset和block的映射关系，其中会涉及block的分配，之前有提到过核心函数就是ext4_map_blocks:

```

static int _ext4_get_block(struct inode *inode, sector_t iblock,
                          struct buffer_head *bh, int flags)
{
    struct ext4_map_blocks map;
    int ret = 0;

    if (ext4_has_inline_data(inode))
        return -ERANGE;

    map.m_lblk = iblock;
    map.m_len = bh->b_size >> inode->i_blkbits;

    ret = ext4_map_blocks(ext4_journal_current_handle(), inode, &map,
                          flags);

    if (ret > 0) {
        map_bh(bh, inode->i_sb, map.m_pblk);
        ext4_update_bh_state(bh, map.m_flags);
        bh->b_size = inode->i_sb->s_blocksize * map.m_len;
        ret = 0;
    }
    return ret;
}

/*
 * The ext4_map_blocks() function tries to look up the requested blocks,
 * and returns if the blocks are already mapped.
 *
 * Otherwise it takes the write lock of the i_data_sem and allocate blocks
 * and store the allocated blocks in the result buffer head and mark it
 * mapped.
 *
 * If file type is extents based, it will call ext4_ext_map_blocks(),
 * Otherwise, call with ext4_ind_map_blocks() to handle indirect mapping
 * based files
 *
 * On success, it returns the number of blocks being mapped or allocated. if
 * create==0 and the blocks are pre-allocated and unwritten, the resulting @map
 * is marked as unwritten. If the create == 1, it will mark @map as mapped.
 *
 * It returns 0 if plain look up failed (blocks have not been allocated), in
 * that case, @map is returned as unmapped but we still do fill map->m_len to
 * indicate the length of a hole starting at map->m_lblk.
 *
 * It returns the error in case of allocation failure.
 */
int ext4_map_blocks(handle_t *handle, struct inode *inode,
                   struct ext4_map_blocks *map, int flags)
{
    struct extent_status es;
    int retval;
    int ret = 0;
#ifdef ES_AGGRESSIVE_TEST
    struct ext4_map_blocks orig_map;

    memcpy(&orig_map, map, sizeof(*map));
#endif

    map->m_flags = 0;
    ext_debug("ext4_map_blocks(): inode %lu, flag %d, max_blocks %u,"
              "logical block %lu\n", inode->i_ino, flags, map->m_len,
              (unsigned long) map->m_lblk);

```

```

/*
 * ext4_map_blocks returns an int, and m_len is an unsigned int
 */
if (unlikely(map->m_len > INT_MAX))//处理len异常
    map->m_len = INT_MAX;

/* We can handle the block number less than EXT_MAX_BLOCKS */
if (unlikely(map->m_lblk >= EXT_MAX_BLOCKS))//处理block number异常
    return -EFSCORRUPTED;

/* Lookup extent status tree firstly */
if (ext4_es_lookup_extent(inode, map->m_lblk, &es)) {
    if (ext4_es_is_written(&es) || ext4_es_is_unwritten(&es)) {
        map->m_pblk = ext4_es_pblock(&es) +
            map->m_lblk - es.es_lblk;
        map->m_flags |= ext4_es_is_written(&es) ?
            EXT4_MAP_MAPPED : EXT4_MAP_UNWRITTEN;
        retval = es.es_len - (map->m_lblk - es.es_lblk);
        if (retval > map->m_len)
            retval = map->m_len;
        map->m_len = retval;
    } else if (ext4_es_is_delayed(&es) || ext4_es_is_hole(&es)) {
        map->m_pblk = 0;
        retval = es.es_len - (map->m_lblk - es.es_lblk);
        if (retval > map->m_len)
            retval = map->m_len;
        map->m_len = retval;
        retval = 0;
    } else {
        BUG_ON(1);
    }
}

#ifdef ES_AGGRESSIVE_TEST
    ext4_map_blocks_es_recheck(handle, inode, map,
                                &orig_map, flags);
#endif

endif
    goto found;
}

/*
 * Try to see if we can get the block without requesting a new
 * file system block.
 */
down_read(&EXT4_I(inode)->i_data_sem);
if (ext4_test_inode_flag(inode, EXT4_INODE_EXTENTS)) {
    retval = ext4_ext_map_blocks(handle, inode, map, flags &
        EXT4_GET_BLOCKS_KEEP_SIZE);
} else {
    retval = ext4_ind_map_blocks(handle, inode, map, flags &
        EXT4_GET_BLOCKS_KEEP_SIZE);
}
if (retval > 0) {
    unsigned int status;

    if (unlikely(retval != map->m_len)) {
        ext4_warning(inode->i_sb,
            "ES len assertion failed for inode "
            "%lu: retval %d != map->m_len %d",
            inode->i_ino, retval, map->m_len);
        WARN_ON(1);
    }

    status = map->m_flags & EXT4_MAP_UNWRITTEN ?
        EXTENT_STATUS_UNWRITTEN : EXTENT_STATUS_WRITTEN;
    if (!(flags & EXT4_GET_BLOCKS_DELALLOC_RESERVE) &&
        !(status & EXTENT_STATUS_WRITTEN) &&
        ext4_find_delalloc_range(inode, map->m_lblk,
            map->m_lblk + map->m_len - 1))
        status |= EXTENT_STATUS_DELAYED;
    ret = ext4_es_insert_extent(inode, map->m_lblk,
        map->m_len, map->m_pblk, status);
    if (ret < 0)
        retval = ret;
}
up_read(&EXT4_I(inode)->i_data_sem);

found:
if (retval > 0 && map->m_flags & EXT4_MAP_MAPPED) {
    ret = check_block_validity(inode, map);
    if (ret != 0)
        return ret;
}

/* If it is only a block(s) look up */
if ((flags & EXT4_GET_BLOCKS_CREATE) == 0)
    return retval;

/*
 * Returns if the blocks have already allocated
 */
/* Note that if blocks have been preallocated
 * ext4_ext_get_block() returns the create = 0
 * with buffer head unmapped.
 */
if (retval > 0 && map->m_flags & EXT4_MAP_MAPPED)
/*
 * If we need to convert extent to unwritten
 * we continue and do the actual work in
 * ext4_ext_map_blocks()
 */
    if (!(flags & EXT4_GET_BLOCKS_CONVERT_UNWRITTEN))
        return retval;

/*
 * Here we clear m_flags because after allocating a new extent,
 * it will be set again.
 */
map->m_flags &= ~EXT4_MAP_FLAGS;

/*
 * New blocks allocate and/or writing to unwritten extent

```



```

* will possibly result in updating i_data, so we take
* the write lock of i_data_sem, and call get_block()
* with create == 1 flag.
*/
down_write(&EXT4_l(inode)->i_data_sem);

/*
 * We need to check for EXT4 here because migrate
 * could have changed the inode type in between
 */
if (ext4_test_inode_flag(inode, EXT4_INODE_EXTENTS)) {
    retval = ext4_ext_map_blocks(handle, inode, map, flags);
} else {
    retval = ext4_ind_map_blocks(handle, inode, map, flags);

    if (retval > 0 && map->m_flags & EXT4_MAP_NEW) {
        /*
         * We allocated new blocks which will result in
         * i_data's format changing. Force the migrate
         * to fail by clearing migrate flags
         */
        ext4_clear_inode_state(inode, EXT4_STATE_EXT_MIGRATE);
    }

    /*
     * Update reserved blocks/metadata blocks after successful
     * block allocation which had been deferred till now. We don't
     * support fallocation for non extent files. So we can update
     * reserve space here.
     */
    if ((retval > 0) &&
        (flags & EXT4_GET_BLOCKS_DELALLOC_RESERVE))
        ext4_da_update_reserve_space(inode, retval, 1);
}

if (retval > 0) {
    unsigned int status;

    if (unlikely(retval != map->m_len)) {
        ext4_warning(inode->i_sb,
            "ES len assertion failed for inode "
            "%lu: retval %d != map->m_len %d",
            inode->i_ino, retval, map->m_len);
        WARN_ON(1);
    }

    /*
     * We have to zeroout blocks before inserting them into extent
     * status tree. Otherwise someone could look them up there and
     * use them before they are really zeroed. We also have to
     * unmap metadata before zeroing as otherwise writeback can
     * overwrite zeros with stale data from block device.
     */
    if (flags & EXT4_GET_BLOCKS_ZERO &&
        map->m_flags & EXT4_MAP_MAPPED &&
        map->m_flags & EXT4_MAP_NEW) {
        ext4_lblk_t i;

        for (i = 0; i < map->m_len; i++) {
            unmap_underlying_metadata(inode->i_sb->s_bdev,
                map->m_pblk + i);
        }
        ret = ext4_issue_zeroout(inode, map->m_lblk,
            map->m_pblk, map->m_len);

        if (ret) {
            retval = ret;
            goto out_sem;
        }
    }

    /*
     * If the extent has been zeroed out, we don't need to update
     * extent status tree.
     */
    if ((flags & EXT4_GET_BLOCKS_PRE_IO) &&
        ext4_es_lookup_extent(inode, map->m_lblk, &es)) {
        if (ext4_es_is_written(&es))
            goto out_sem;
    }

    status = map->m_flags & EXT4_MAP_UNWRITTEN ?
        EXTENT_STATUS_UNWRITTEN : EXTENT_STATUS_WRITTEN;
    if (!(flags & EXT4_GET_BLOCKS_DELALLOC_RESERVE) &&
        !(status & EXTENT_STATUS_WRITTEN) &&
        ext4_find_delalloc_range(inode, map->m_lblk,
            map->m_pblk + map->m_len - 1))
        status |= EXTENT_STATUS_DELAYED;
    ret = ext4_es_insert_extent(inode, map->m_lblk, map->m_len,
        map->m_pblk, status);

    if (ret < 0) {
        retval = ret;
        goto out_sem;
    }
}

out_sem:
up_write(&EXT4_l(inode)->i_data_sem);
if (retval > 0 && map->m_flags & EXT4_MAP_MAPPED) {
    ret = check_block_validity(inode, map);
    if (ret != 0)
        return ret;
}

/*
 * Inodes with freshly allocated blocks where contents will be
 * visible after transaction commit must be on transaction's
 * ordered data list.
 */
if (map->m_flags & EXT4_MAP_NEW &&
    !(map->m_flags & EXT4_MAP_UNWRITTEN) &&
    !(flags & EXT4_GET_BLOCKS_ZERO) &&
    !IS_NOQUOTA(inode) &&
    ext4_should_order_data(inode)) {
    if (flags & EXT4_GET_BLOCKS_IO_SUBMIT)

```

```
        ret = ext4_jbd2_inode_add_wait(handle, inode);
    else
        ret = ext4_jbd2_inode_add_write(handle, inode);
    if (ret)
        return ret;
    }
}
return retval;
}
```