

# Async/sync/fsync/fdatasync

2018年5月11日 14:13

使用区别？

实现差异？

File open flags:

O\_SYNC

Write operations on the file will complete according to the requirements of synchronized I/O *file* integrity completion (by contrast with the synchronized I/O *data* integrity completion provided by **O\_DSYNC**.)

By the time [write\(2\)](#) (or similar) returns, the output data and associated file metadata have been transferred to the underlying hardware (i.e., as though each [write\(2\)](#) was followed by a call to [fsync\(2\)](#)).

O\_DSYNC

Write operations on the file will complete according to the requirements of synchronized I/O *data* integrity completion.

By the time [write\(2\)](#) (and similar) return, the output data has been transferred to the underlying hardware, along with any file metadata that would be required to retrieve that data (i.e., as though each [write\(2\)](#) was followed by a call to [fdatasync\(2\)](#)).

O\_ASYNC

Enable signal-driven I/O: generate a signal (**SIGIO** by default, but this can be changed via [fcntl\(2\)](#)) when input or output becomes possible on this file descriptor. This feature is available only for terminals, pseudoterminals, sockets, and (since Linux 2.6) pipes and FIFOs. See [fcntl\(2\)](#) for further details.

系统调用：

Sync: 唤醒disk flush线程，将dirty数据提交到io队列，总是返回成功

SYSCALL\_DEFINE0(sync)

```
{
    int nowait = 0, wait = 1;

    wakeup_flusher_threads(0, WB_REASON_SYNC);
    iterate_supers(sync_inodes_one_sb, NULL);
    iterate_supers(sync_fs_one_sb, &nowait);
    iterate_supers(sync_fs_one_sb, &wait);
    iterate_bdevs(fdatawrite_one_bdev, NULL);
    iterate_bdevs(fdatawait_one_bdev, NULL);
    if (unlikely(laptop_mode))
        laptop_sync_completion();
    return 0;
}
```

Syncfs: 根据传入的参数wait/nowait，返回可能失败

SYSCALL\_DEFINE1(syncfs, int, fd)

```
{
    struct fd f = fdget(fd);
    struct super_block *sb;
    int ret;

    if (!f.file)
        return -EBADF;
```

```

        sb = f.file->f_path.dentry->d_sb;

        down_read(&sb->s_umount);
        ret = sync_filesystem(sb);
        up_read(&sb->s_umount);

        fdput(f);
        return ret;
}
int sync_filesystem(struct super_block *sb)
{
    int ret;

    /*
     * We need to be protected against the filesystem going from
     * r/o to r/w or vice versa.
     */
    WARN_ON(!rwsem_is_locked(&sb->s_umount));

    /*
     * No point in syncing out anything if the filesystem is read-only.
     */
    if (sb->s_flags & MS_RDONLY)
        return 0;

    ret = __sync_filesystem(sb, 0);
    if (ret < 0)
        return ret;
    return __sync_filesystem(sb, 1);
}
static int __sync_filesystem(struct super_block *sb, int wait)
{
    if (wait)
        sync_inodes_sb(sb);
    else
        writeback_inodes_sb(sb, WB_REASON_SYNC);

    if (sb->s_op->sync_fs)
        sb->s_op->sync_fs(sb, wait);
    return __sync_blockdev(sb->s_bdev, wait);
}

```

Fsync: 等待文件io完成落盘后返回

```

SYSCALL_DEFINE1(fsync, unsigned int, fd)
{
    return do_fsync(fd, 0);
}

```

fdatsync: 等待文件的数据io完成落盘后返回

```

SYSCALL_DEFINE1(fdatsync, unsigned int, fd)
{
    return do_fsync(fd, 1);
}

```

```

int vfs_fsync(struct file *file, int datasync)
{
    return vfs_fsync_range(file, 0, LLONG_MAX, datasync);
}

```

```
static int do_fsync(unsigned int fd, int datasync)
{
    struct fd f = fdget(fd);
    int ret = -EBADF;

    if (f.file) {
        ret = vfs_fsync(f.file, datasync);
        fdput(f);
    }
    return ret;
}
```

sync\_file\_range: 刷文件的一部分

Since glibc 2.2.2, the Linux prototype for **sync()** is as listed above, following the various standards. In glibc 2.2.1 and earlier, it was "int sync(void)", and **sync()** always returned 0. According to the standard specification (e.g., POSIX 1-2001), **sync()** schedules the writes, but may return before the actual writing is done. However Linux waits for I/O completions, and thus **sync()** or **syncfs()** provide the same guarantees as fsync called on every file in the system or filesystem respectively.

## BUGS

[top](#)

Before version 1.3.20 Linux did not wait for I/O to complete before returning.

来自 <<http://man7.org/linux/man-pages/man2/sync.2.html>>